

# MIE1628: BIG DATA SCIENCE

## ASSIGNMENT 2: Spark

Bowen Xu 1006411786

**Due on: November 11th**

### Section 1: Machine Learning and Time Series Prediction (25 points)

Given Data: APPLE daily close price from 2016-2017 (2 years inclusive)

Download it from Yahoo finance: <https://finance.yahoo.com/quote/AAPL/history?p=AAPL>

**Q1:** Use apple close price data and create two lag features, lag1 and lag2. Lag1 should push the close price back by one day and lag2 should push the price back by two days. Show your code and data in the dataframe.

```
import org.apache.spark.sql.expressions.Window
import org.apache.spark.sql.functions._
import org.apache.spark._
import org.apache.spark.ml.linalg.{Vector, Vectors}
import org.apache.spark.ml.feature.{StringIndexerModel, VectorAssembler}
import org.apache.spark.ml.evaluation.RegressionEvaluator
import org.apache.spark.ml.regression.LinearRegression

import org.apache.spark.sql.expressions.Window
import org.apache.spark.sql.functions._
import org.apache.spark._
import org.apache.spark.ml.linalg.{Vector, Vectors}
import org.apache.spark.ml.feature.{StringIndexerModel, VectorAssembler}
import org.apache.spark.ml.evaluation.RegressionEvaluator
import org.apache.spark.ml.regression.LinearRegression

//Q1: Lag1 should push the close price back by one day and lag2 should push the price back by two days.
val df3 = spark.sql("select date, close from aapl_csv where YEAR(date) between 2016 and 2017")
val partitionWindow = Window.orderBy("date")
val lag1 = lag("Close",1,0).over(partitionWindow)
val lag2 = lag("Close",2,0).over(partitionWindow)
val data = df3.withColumn("lag1",lag1).withColumn("lag2",lag2).select("*")
data.show(10)

+ (1) Spark Jobs
+ df3: org.apache.spark.sql.DataFrame = [date: timestamp, close: double]
+ data: org.apache.spark.sql.DataFrame = [date: timestamp, close: double ... 2 more fields]
+-----+
|      date|      close|      lag1|      lag2|
+-----+
|2016-01-04 00:00:00|105.349998|         0.0|         0.0|
|2016-01-05 00:00:00|102.709999|105.349998|         0.0|
|2016-01-06 00:00:00|100.699997|102.709999|105.349998|
|2016-01-07 00:00:00| 96.449997|100.699997|102.709999|
|2016-01-08 00:00:00| 96.959999| 96.449997|100.699997|
|2016-01-11 00:00:00| 98.529999| 96.959999| 96.449997|
|2016-01-12 00:00:00| 99.959999| 98.529999| 96.959999|
|2016-01-13 00:00:00| 97.389999| 99.959999| 98.529999|
|2016-01-14 00:00:00| 99.519997| 97.389999| 99.959999|
|2016-01-15 00:00:00| 97.129997| 99.519997| 97.389999|
+-----+
only showing top 10 rows
```

**Q2:** Split dataframe into train and test (0.7/0.3) and train your model use linear regression on 70% of your data and test with the other 30 percent. Show your Code.

```
//Q2: Split dataframe into train and test (0.7/0.3) and train your model use linear regression on 70% of your data and test with the other 30 percent. Show your Code.(RandomSplit)
val data1 = data.select("lag1","lag2","Close")
val assembler = new VectorAssembler().setInputCols(Array("lag1","lag2")).setOutputCol("features")
val output = assembler.transform(data1)
val training = output.select("features","Close").toDF("features","label")
val Array(train_data,test_data) = training.randomSplit(Array(0.7, 0.3))
val regressor = new LinearRegression()
val model = regressor.fit(train_data)
val pred_results = model.transform(test_data)

+ (5) Spark Jobs
+ data1: org.apache.spark.sql.DataFrame = [lag1: double, lag2: double ... 1 more fields]
+ output: org.apache.spark.sql.DataFrame = [lag1: double, lag2: double ... 2 more fields]
+ training: org.apache.spark.sql.DataFrame = [features: udt, label: double]
+ train_data: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [features: udt, label: double]
+ test_data: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [features: udt, label: double]
+ pred_results: org.apache.spark.sql.DataFrame = [features: udt, label: double ... 1 more fields]
data1: org.apache.spark.sql.DataFrame = [lag1: double, lag2: double ... 1 more field]
assembler: org.apache.spark.ml.feature.VectorAssembler = vecAssembler_4942c190f321
output: org.apache.spark.sql.DataFrame = [lag1: double, lag2: double ... 2 more fields]
training: org.apache.spark.sql.DataFrame = [features: vector, label: double]
train_data: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [features: vector, label: double]
test_data: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [features: vector, label: double]
regressor: org.apache.spark.ml.regression.LinearRegression = linReg_72adcc1bc72e
model: org.apache.spark.ml.regression.LinearRegressionModel = linReg_72adcc1bc72e
pred_results: org.apache.spark.sql.DataFrame = [features: vector, label: double ... 1 more field]
```

**Q3:** Create a prediction column and show your prediction. Show your code and dataframe.

```
//Q3: Create a prediction column and show your prediction. Show your code and dataframe.
pred_results.show(10)
```

```
▶ (1) Spark Jobs

+-----+-----+-----+
| features | label | prediction |
+-----+-----+-----+
|[92.040001,93.400...| 93.589996|93.95258962754873|
|[92.790001,92.720...| 93.419998| 94.866733364036|
|[93.419998,92.790...| 92.510002|95.54498008992842|
|[93.879997,90.519...| 93.489998|96.36469532744464|
|[ 94.190002,95.18]| 93.239998|96.05180117201732|
|[ 96.43,97.339996]| 94.480003|98.19618648950942|
|[96.669998,97.339...|102.949997| 98.458293192755|
|[96.790001,96.660...| 96.300003|98.68441211777402|
|[96.870003,97.419...| 98.790001|98.66553864392638|
|[96.959999,96.449...| 98.529999|98.89942885627954|
+-----+-----+-----+

only showing top 10 rows
```

**Q4:** Evaluate your model with evaluation metrics (RMSE), show your code and print your result.

```
//Q4: Evaluate your model with evaluation metrics (RMSE), show your code and print your result.
val evaluator= new RegressionEvaluator().setLabelCol("label").setPredictionCol("prediction").setMetricName("rmse")
val rmse = evaluator.evaluate(pred_results)
println("Root Mean Squared Error (RMSE) on test data: ",rmse)
```

```
▶ (1) Spark Jobs

(Root Mean Squared Error (RMSE) on test data: ,2.5604472742348228)
evaluator: org.apache.spark.ml.evaluation.RegresionEvaluator = regEval_73317de9f99c
rmse: Double = 2.5604472742348228
```

**Question5 - Bonus:** Play around with features(Only) and try to see if you can get better result(split 0.7/0.3 and linear regression), if you can get top 5 result in the class, you will be reward with 100 percent on this assignment.

*Features which can get lowest RMSE are "Open", "lag2", "Adj Close", results are:*

BONUS

```
import org.apache.spark.sql.expressions.Window
import org.apache.spark.sql.functions._
import org.apache.spark._
import org.apache.spark.ml.linalg.{Vector, Vectors}
import org.apache.spark.ml.feature.{StringIndexerModel, VectorAssembler }
import org.apache.spark.ml.evaluation.RegresionEvaluator
import org.apache.spark.ml.regression.LinearRegression

val df4 = spark.sql("select * from aspl_csv where YEAR(date) between 2016 and 2017")
val partitionWindow1 = Window.orderBy("Date")
val lag11 = lag("Close",1,0).over(partitionWindow1)
val lag22 = lag("Close",2,0).over(partitionWindow1)
val data2 = df4.withColumn("lag1",lag11).withColumn("lag22",lag22).select(".*")

val data3 = data2.select("Open","lag2","Adj Close","Close")
val assembler1 = new VectorAssembler().setInputCols(Array("Open","lag2","Adj Close")).setOutputCol("features")
val output1 = assembler1.transform(data3)
val training1 = output1.select("features","Close").toDF("features","label")
val Array(train_data1,test_data1) = training1.randomSplit(Array(0.7, 0.3))
val regressor1 = new LinearRegression()
val model1 = regressor1.fit(train_data1)
val pred_results1 = model1.transform(test_data1)

val evaluator1= new RegressionEvaluator().setLabelCol("label").setPredictionCol("prediction").setMetricName("rmse")
val rmse1 = evaluator1.evaluate(pred_results1)
println("Root Mean Squared Error (RMSE) on test data: ",rmse1)
```

```
▶ (6) Spark Jobs

df4: org.apache.spark.sql.DataFrame = [Date: timestamp, Open: double ... 5 more fields]
data2: org.apache.spark.sql.DataFrame = [Date: timestamp, Open: double ... 7 more fields]
data3: org.apache.spark.sql.DataFrame = [Open: double, lag2: double ... 2 more fields]
output1: org.apache.spark.sql.DataFrame = [Open: double, lag2: double ... 3 more fields]
training1: org.apache.spark.sql.DataFrame = [features: udt, label: double]
train_data1: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [features: udt, label: double]
test_data1: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [features: udt, label: double]
pred_results1: org.apache.spark.sql.DataFrame = [features: udt, label: double ... 1 more fields]

(Root Mean Squared Error (RMSE) on test data: ,0.37028545192433016)
import org.apache.spark.sql.expressions.Window
import org.apache.spark.sql.functions._
import org.apache.spark._
import org.apache.spark.ml.linalg.{Vector, Vectors}
import org.apache.spark.ml.feature.{StringIndexerModel, VectorAssembler}
import org.apache.spark.ml.evaluation.RegresionEvaluator
import org.apache.spark.ml.regression.LinearRegression

df4: org.apache.spark.sql.DataFrame = [Date: timestamp, Open: double ... 5 more fields]
partitionWindow1: org.apache.spark.sql.expressions.WindowSpec = org.apache.spark.sql.expressions.WindowSpec@1687d734
lag11: org.apache.spark.sql.Column = lag(Close, 1, 0) OVER (ORDER BY Date ASC NULLS FIRST unspecifiedFrames())
lag22: org.apache.spark.sql.Column = lag(Close, 2, 0) OVER (ORDER BY Date ASC NULLS FIRST unspecifiedFrames())
data2: org.apache.spark.sql.DataFrame = [Date: timestamp, Open: double ... 7 more fields]
data3: org.apache.spark.sql.DataFrame = [Open: double, lag2: double ... 2 more fields]
assembler1: org.apache.spark.ml.feature.VectorAssembler = vecAssembler_312df00c8563
output1: org.apache.spark.sql.DataFrame = [Open: double, lag2: double ... 3 more fields]
training1: org.apache.spark.sql.DataFrame = [features: vector, label: double]
train_data1: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [features: vector, label: double]
test_data1: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [features: vector, label: double]
regressor1: org.apache.spark.ml.regression.LinearRegression = linReg_c9412f5fefac
model1: org.apache.spark.ml.regression.LinearRegressionModel = linReg_c9412f5fefac
pred_results1: org.apache.spark.sql.DataFrame = [features: vector, label: double ... 1 more field]
evaluator1: org.apache.spark.ml.evaluation.RegresionEvaluator = regEval_adai05753ee6
rmse1: Double = 0.37028545192433016
```

## Section 2: PySpark/Scala syntax

### Question 1 (15 points)

1. Select two columns - games and seasons - and add a column with total goals (sum of home and away goals). Suggestion: use `df.withColumn()` function.

```
import org.apache.spark.sql.functions._
val df = spark.table("game_csv")
val table1 = df.withColumn("total_goals", $"home_goals" + $"away_goals").select("season", "game_id", "total_goals")
table1.show(10)
```

▶ (1) Spark Jobs

▶ df: org.apache.spark.sql.DataFrame = [game\_id: integer, season: integer ... 14 more fields]

▶ table1: org.apache.spark.sql.DataFrame = [season: integer, game\_id: integer ... 1 more fields]

```
+-----+-----+-----+
| season| game_id|total_goals|
+-----+-----+-----+
|20112012|2011030221|7|
|20112012|2011030222|5|
|20112012|2011030223|7|
|20112012|2011030224|6|
|20112012|2011030225|4|
|20112012|2011030411|3|
|20112012|2011030412|3|
|20112012|2011030413|4|
|20112012|2011030414|4|
|20112012|2011030415|3|
+-----+-----+-----+
only showing top 10 rows
```

```
import org.apache.spark.sql.functions._
df: org.apache.spark.sql.DataFrame = [game_id: int, season: int ... 14 more fields]
table1: org.apache.spark.sql.DataFrame = [season: int, game_id: int ... 1 more field]
```

2. Organize records in ascending order (by season).

```
table1.orderBy($"season").show(10)
```

▶ (1) Spark Jobs

```
+-----+-----+-----+
| season| game_id|total_goals|
+-----+-----+-----+
|20102011|2010030311|7|
|20102011|2010030244|7|
|20102011|2010030312|11|
|20102011|2010030313|2|
|20102011|2010030314|8|
|20102011|2010030315|4|
|20102011|2010030316|9|
|20102011|2010030317|1|
|20102011|2010030241|3|
|20102011|2010030242|3|
+-----+-----+-----+
only showing top 10 rows
```

3. Add a column with an average, min and max total score for each season. Suggestion: use Window function.

```
import org.apache.spark.sql.expressions.Window
val partitionwindow = Window.partitionBy($"season")
val table2 = table1.select($"*", avg("total_goals") over(partitionwindow) as "avg_total_goals",
                          min("total_goals") over(partitionwindow) as "min_total_goals",
                          max("total_goals") over(partitionwindow) as "max_total_goals")
table2.show(10)
```

▶ (1) Spark Jobs

```
▶ table2: org.apache.spark.sql.DataFrame = [season: integer, game_id: integer ... 4 more fields]
+-----+-----+-----+-----+-----+
| season| game_id|total_goals| avg_total_goals|min_total_goals|max_total_goals|
+-----+-----+-----+-----+-----+
|20112012|2011030221|7|5.427051671732523|1|17|
|20112012|2011030222|5|5.427051671732523|1|17|
|20112012|2011030223|7|5.427051671732523|1|17|
|20112012|2011030224|6|5.427051671732523|1|17|
|20112012|2011030225|4|5.427051671732523|1|17|
|20112012|2011030411|3|5.427051671732523|1|17|
|20112012|2011030412|3|5.427051671732523|1|17|
|20112012|2011030413|4|5.427051671732523|1|17|
|20112012|2011030414|4|5.427051671732523|1|17|
|20112012|2011030415|3|5.427051671732523|1|17|
+-----+-----+-----+-----+-----+
only showing top 10 rows
```

```
import org.apache.spark.sql.expressions.Window
partitionwindow: org.apache.spark.sql.expressions.WindowSpec = org.apache.spark.sql.expressions.WindowSpec@62a3f765
table2: org.apache.spark.sql.DataFrame = [season: int, game_id: int ... 4 more fields]
```

4. Add a column that finds a difference between each game's total score and average for that season. Suggestion: use Window function.

```
val table3 = table2.withColumn("diff_total_goals", $"total_goals" - $"avg_total_goals")
table3.show(10)
```

▶ (4) Spark Jobs

table3: org.apache.spark.sql.DataFrame = [season: integer, game\_id: integer ... 5 more fields]

season	game_id	total_goals	avg_total_goals	min_total_goals	max_total_goals	diff_total_goals
20112012	2011030221	7	5.427051671732523	1	17	1.5729483282674774
20112012	2011030222	5	5.427051671732523	1	17	-0.42705167173252256
20112012	2011030223	7	5.427051671732523	1	17	1.5729483282674774
20112012	2011030224	6	5.427051671732523	1	17	0.5729483282674774
20112012	2011030225	4	5.427051671732523	1	17	-1.4270516717325226
20112012	2011030411	3	5.427051671732523	1	17	-2.4270516717325226
20112012	2011030412	3	5.427051671732523	1	17	-2.4270516717325226
20112012	2011030413	4	5.427051671732523	1	17	-1.4270516717325226
20112012	2011030414	4	5.427051671732523	1	17	-1.4270516717325226
20112012	2011030415	3	5.427051671732523	1	17	-2.4270516717325226

only showing top 10 rows

table3: org.apache.spark.sql.DataFrame = [season: int, game\_id: int ... 5 more fields]

5. Print top 10 records.

```
table3.orderBy($"diff_total_goals".desc).show(10)
```

▶ (1) Spark Jobs

season	game_id	total_goals	avg_total_goals	min_total_goals	max_total_goals	diff_total_goals
20112012	2011020128	17	5.427051671732523	1	17	11.572948328267477
20162017	2016020661	15	5.506454062262718	1	15	9.49354593773728
20102011	2010020271	15	5.586808188021228	1	15	9.413191811978772
20182019	2018020420	15	6.000736377025037	1	15	8.999263622974963
20182019	2018020912	15	6.000736377025037	1	15	8.999263622974963
20162017	2016020434	14	5.506454062262718	1	15	8.49354593773728
20102011	2010020808	14	5.586808188021228	1	15	8.413191811978772
20122013	2012020306	13	5.398263027295285	1	13	7.601736972704715
20112012	2011020183	13	5.427051671732523	1	17	7.572948328267477
20112012	2011020398	13	5.427051671732523	1	17	7.572948328267477

only showing top 10 rows

## Question 2 (10 points)

1. List all team names (teamName) for teams that played as away team at TD Garden during seasons 2012-2013 and 2013-2014.

```
val table4 = spark.table("game_csv").filter(($"venue" === "TD Garden") && (($"season" === "20122013") || ($"season" === "20132014"))).select("away_team_id")
val df2 = spark.table("team_info_csv").select("teamname", "team_id")
val table_join = df2.join(table4, table4.col("away_team_id") === df2.col("team_id"))
val table_name = table_join.select("teamname")
table_name.show()
```

▶ (1) Spark Jobs

table4: org.apache.spark.sql.DataFrame = [away\_team\_id: integer]

df2: org.apache.spark.sql.DataFrame = [teamname: string, team\_id: integer]

table\_join: org.apache.spark.sql.DataFrame = [teamname: string, team\_id: integer ... 1 more fields]

table\_name: org.apache.spark.sql.DataFrame = [teamname: string]

teamname
Rangers
Rangers
Rangers
Penguins
Penguins
Maple Leafs
Maple Leafs
Maple Leafs
Blackhawks
Blackhawks
Blackhawks
Red Wings
Red Wings
Red Wings
Canadiens
Canadiens
Canadiens
Canadiens
Canadiens
Hurricanes

only showing top 20 rows

table4: org.apache.spark.sql.DataFrame = [away\_team\_id: int]

df2: org.apache.spark.sql.DataFrame = [teamname: string, team\_id: int]

table\_join: org.apache.spark.sql.DataFrame = [teamname: string, team\_id: int ... 1 more field]

table\_name: org.apache.spark.sql.DataFrame = [teamname: string]

## 2. How many unique teams are on the list?

```
table_name.distinct().count()
```

```
↳ (1) Spark Jobs  
res37: Long = 29
```

### Question 3 (additional 15 points)

(Bonus Question (doesn't have to be completed to get a full mark for this assignment, but if you complete it you will get additional 15 points))

- Create a function that when input a number  $n$  returns a list of prime numbers between 1 and  $n$ .
- Test your function with number 17.

```
def getPrimeList(n: Int) = {  
  require(n >= 2)  
  val oddList = 3 to n by 2 toList  
  def pn(oddList: List[Int], primeList: List[Int]): List[Int] = oddList match {  
    case Nil => primeList  
    case _ if primeList.exists(oddList.head % _ == 0) => pn(oddList.tail, primeList)  
    case _ => pn(oddList.tail, oddList.head :: primeList)  
  }  
  pn(oddList, List(2)).reverse  
}
```

```
warning: there was one feature warning; re-run with -feature for details  
getPrimeList: (n: Int)List[Int]
```

```
getPrimeList(17)
```

```
res2: List[Int] = List(2, 3, 5, 7, 11, 13, 17)
```